

Lecture 1

Marc Kaufmann

9/10/2019

This introduction is essentially Chapter 2 from Kieran Healy’s book Data Visualization, which you should read as part of Assignment 2 and to find your way around RStudio. The reason I start with this (and continue for one or two more weeks) is that visualization is one of the more fun parts of data analysis, as well as one you should go through – so it makes for a great starting point.

We will use exclusively R Markdown files. In this case, you should simply open the file ‘lecture1-to-fill.Rmd’ from the repository and you are set. Whenever you create a new file, you should choose to make an R Markdown file.

Things to Know about R

There are 4 things to bear in mind about R:

1. Everything has a name
2. Everything is an object
3. You do things using functions
4. Functions come in packages

Everything has a name

Everything that you use in R has a name: variables (including datasets), functions, or special reserved words. That is the way you talk about them. Here are some examples:

```
3 # Numbers are called by their number, arithmetic operations by their usual symbol

c # There are some pre-defined variables

str() # There are also pre-defined functions
```

Some Notes:

- ‘#’ is the Comment (EXERCISE) sign. It tells R that everything that follows is a comment and R should ignore it. Comments are there to help explain parts of the code that need additional documentation. We will overuse them initially, but reduce this as the course proceeds.
- c, “c”, and “C” are not the same things
- Naming conventions: When you name variables and functions, you should use snake case.

Exercise: For each of the lines in the following code chunk, write in a comment next to it what it returns. I completed the first example for you.

```
3          # -> 3
7 %% 3     # -> 1
False      # -> Error: Object False not found
'pi'       # -> 3.141593
FALSE      # -> FALSE
```

Exercise: Snake case (or snake_case) is the practice of writing compound words or phrases in which the elements are separated with one underscore character (__) and no spaces, with each element’s initial letter usually lowercased within the compound and the first letter either upper- or lowercase—as in “foo_bar” and “Hello_world”

Exercise: Identifiers can be a combination of letters, digits, period (.) and underscore (_). It must start with a letter or a period. If it starts with a period, it cannot be followed by a digit. Reserved words in R cannot be used as identifiers. Examples of not valid names: `.66value`, `_value`, `c`

Everything is an object

There are built-in objects and objects you import or create. Most importantly, you assign values to objects you create with the `<-` operator: a `<` ('less than') followed by a `-` ('minus' or 'dash').

```
marcs_new_object <- 6
marcs_new_object
```

One important function to create objects is the `c()` function, which combines several items into one objects:

```
marcs_new_combined_object <- c(1, 2, 3, 4, 5)
marcs_new_combined_object
```

In principle, you can use `=` to assign values, but this is very non-idiomatic R. I will subtract points for using `=`. Use `<-`.

Class Exercise: Let's start doing something with data. How many of the following programming languages have you *heard* of?

```
# Indent the list
list_of_programming_languages <- c(
  "R",
  "SQL",
  "Racket",
  "Lisp",
  "JavaScript",
  "ECMAScript",
  "bash",
  "C",
  "Perl",
  "Logo",
  "Scratch"
)

languages_heard_of <- c(6)
```

You do things using functions

A function is a thing that does things to things.

Me, paraphrasing Cosma Shalizi

A function is a special object that you can call, such as the function *mean*. It is an object:

```
mean

## function (x, ...)
## UseMethod("mean")
## <bytecode: 0x0000000014615f40>
## <environment: namespace:base>
```

When we *call* it, we put parentheses at the end, and we tell it what to perform its action on:

```
mean(languages_heard_of)

## [1] 6
```

Exercise: What happens if you call the function *mean* without any arguments, i.e. *mean()*?

We get the following error: Error in mean.default() : argument “x” is missing, with no default

Functions come in packages

As you may start to realize, lots of functions already exist, such as *mean()* and *sd()* and *c()*. There are many more functions that were written by other people and bundled into packages so that you can use them. Many packages come bundled with R from the start (base R), while you can install others on your system via ‘install.packages()’ and use via ‘library()’. We will soon use the ‘ggplot()’ function. Try using it now:

```
ggplot()
```

You will get an error. We first need to import the library:

Yes I got an error

Error: object ‘ggplot’ not found

```
#  
library(ggplot2)  
ggplot()
```

This tells R to go and get all the objects defined in the library called *ggplot2*, which includes *ggplot*. This gives us our first data (non-)visualization: a beautifully empty plot.

Think of it as the equivalent of recipe books: you can either experiment and concoct your own recipes, which is fine if you are a cook. Alternatively, you get recipe books from great chefs and follow their instructions. Depending on the type of recipe, you still need a lot of background knowledge – what does it mean to boil, fry, blanch, roast – but you don’t have to figure out everything.

How to figure out what is what

Suppose you have an object x that you don't know what it is. You can do a few things to find out:

```
x <- c(1, 3.0, 2.9)
str(x)

##  num [1:3] 1 3 2.9
attributes(x)
```

```
## NULL
```

For built-in functions, you can also ask for help or bring up the documentation with 'help()' or '?' or '??':


```
# ?c
```

Exercise: What do 'class()' and 'str()' do? Use 'help' (or '?'). Don't spend too much time reading the docs. Which description do you find more helpful?

Class() converts any object to a class() (for an object oriented style of programming) so that constructors and others stuff can be done that is usually associated with OOP

Make Your First Figure

```
library(gapminder)
p = ggplot(data = gapminder, mapping = aes(x = gdpPercap, y = lifeExp))
p + geom_point()
```



As you can see, some datasets are conveniently bundled for you as libraries.

Assignment 2

This assignment is due before the start of class 2. Commit your knitted assignment2.html to your git and push it to your repository regularly as you make progress.

- Only your work until _____ (EXERCISE) will be considered.
- If you know that you might miss the deadline, you have to email me in advance and I will tell you whether you can get an extension. No extension will be granted less than 24 hours prior to the deadline – you should have started working on it.
- If you miss the deadline (even by 1 minute) you lose 25% irrespective of technical issues such as ‘the internet went down’. It wasn’t down the whole week, and you knew a deadline was coming. In a business setting, you would probably lose more than 25%. Repeat offenders will face an increasingly large penalty.
- If you struggle *answering* the assignment, you should submit what you tried, and send me a message that you struggled. Much of the grade initially is on trying out things, even if it doesn’t work out, so you should submit.

Part 1

Read chapter 2 of Data Visualization. Make sure you know where the console is, the editor, and what code chunks are in RStudio.

Part 2

Complete all the exercises in the lectures notes of lecture 1. Put your answer to an exercise right after the question of the exercise in the R Markdown file. When done, knit the document, commit the changes with the commit message “Part 1 of assignment 2” and push them to your GitHub repository.

Part 3

Fill in the holes in the lecture notes of lecture 1, if you didn’t complete it during class. When done, knit the document, commit the changes with the commit message “Part 3 of assignment 2” and push them to your GitHub repository.

Knit the slides for class 1 to html and/or pdf, commit the .Rmd to your repository – but not the html or pdf – and push. If the pdf is hard to get working (due to weird error messages), see whether you can find out *what* the problem is (Stackoverflow, Discourse...) without trying to solve it yet. Pdf and Latex issues can be... interesting shall we say.

Note: If you can’t figure out how to use git, post a question on the discourse forum and send me the html file by email.

Part 4

Track as many error messages as possible that you made during this week (stop once you get to 4). Put the code that caused the error in the chunk below, and copy the error message as a comment below it. I provide an example.

```
# Example: "4" - "3" leads to the error 'Error in "4" - "3" : non-numeric argument to binary operator'
# Thus you write:
"4" - "3"
# -> Error in "4" - "3" : non-numeric argument to binary operator

# Your examples below

#1. When trying to knit as PDF I was getting the following error despite having installed MiKTeX.
#This is a common issue according to their website and can be solved by checking the always
#install missing packages on the fly as yes in MiKTeX settings.
```

```

#output file: lecture1-slides.knit.md

#! Sorry, but C:\Users\faaez\AppData\Local\Programs\MIKTeX~1.9\miktex\bin\x64\xelatex.exe
#did not succeed.

#! The log file hopefully contains the information to get MiKTeX going again:

#! C:\Users\faaez\AppData\Local\MiKTeX\2.9\miktex\log\xelatex.log

#Error: Failed to compile lecture1-slides.tex. See https://yihui.name/tinytex/r/#debugging
#for debugging tips. See lecture1-slides.log for more info.
#Execution halted

?x

#2. Not an error but still was doing something different to what I wanted -
#I was trying to do ?x for my array/list named x and it kept showing me details
#about positional aesthetics that use x and y for mapping.
#Using str() or attribute() showed me what I wanted.

install.package("gapminder")
#3. Kept getting an error because I was doing install.package(gapminder)
#from within the RMD file - learned that it is not advisable to do that and did it using my
#console. Error in contrib.url(repos, "source") :
# trying to use CRAN without setting a mirror
#Calls: <Anonymous> ... withVisible -> eval -> eval -> install.packages ->
#contrib.url

p = ggplot(gdpPercap, lifeExp)
#4. Was trying to use ggplot without defining data and mapping and
#directly inputting the x and y values. Got fixed when I correctly
#defined the data source and mapping using aes -
#corrected code: p = ggplot(data = gapminder, mapping = aes(x = gdpPercap, y = lifeExp))
#Error:
#Error in ggplot(gdpPercap, lifeExp) : object 'gdpPercap' not found
#Calls: <Anonymous> ... withCallingHandlers -> withVisible -> eval ->
#eval -> ggplot
#Execution halted

#Unable to Knit slides
#5. No LaTeX installation detected (LaTeX is required to create PDF output).
#You should install a LaTeX distribution for your platform:
#https://www.latex-project.org/get/

# If you are not sure, you may install TinyTeX in R: tinytex::install_tinytex()

#Otherwise consider MiKTeX on Windows - http://miktex.org

#MacTeX on macOS - https://tug.org/mactex/
#(NOTE: Download with Safari rather than Chrome _strongly_ recommended)

#Linux: Use system package manager

```

Post one of these 4 errors on discourse.

Part 5

Try to figure out two of the following and post your answers on Discourse for others to read:

- Identify elements from my *eureka_or_bust* example in the slides and see whether you can figure out what the different elements mean to R.
- What does *eval=FALSE* mean in part 4? Figure this out by replacing it by *eval=TRUE* and seeing what you get.
- What is the Tidyverse?
- Who is Hadley Wickham?

Resources for this lecture