

Assignment for Lecture 4

R Markdown

mutate()

```
library(nycflights13)
library(tidyverse)
```

```
## -- Attaching packages -----
## v ggplot2 3.2.1    v purrr  0.3.2
## v tibble  2.1.3    v dplyr  0.8.3
## v tidyr   0.8.3    v stringr 1.4.0
## v readr   1.3.1    v forcats 0.4.0

## -- Conflicts ----- tidy
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(ggpubr)
```

```
## Loading required package: magrittr
```

```
##
```

```
## Attaching package: 'magrittr'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
##      set_names
```

```
## The following object is masked from 'package:tidyr':
```

```
##
```

```
##      extract
```

```
# select(flights, dep_time, arr_time, air_time)
```

Let's stare at the columns to see what we can choose from View(flights)

Narrow the tibble to see what mutate() is doing

```
(flights_small <- select(flights,
  year:day,
  ends_with("delay"),
  distance,
  air_time))
```

```
## # A tibble: 336,776 x 7
```

```
##   year month   day dep_delay arr_delay distance air_time
##   <int> <int> <int>   <dbl>   <dbl>   <dbl>   <dbl>
## 1  2013     1     1         2        11    1400     227
## 2  2013     1     1         4        20    1416     227
## 3  2013     1     1         2        33    1089     160
## 4  2013     1     1        -1       -18    1576     183
## 5  2013     1     1        -6       -25     762     116
## 6  2013     1     1        -4        12     719     150
## 7  2013     1     1        -5        19    1065     158
## 8  2013     1     1        -3       -14     229      53
## 9  2013     1     1        -3        -8     944     140
```

```
## 10 2013      1      1      -2          8      733      138
## # ... with 336,766 more rows
```

```
mutate(flights_small,
  catchup = dep_delay - arr_delay,
  speed_miles = (distance/air_time) * 60
)
```

```
## # A tibble: 336,776 x 9
##   year month   day dep_delay arr_delay distance air_time catchup
##   <int> <int> <int>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1  2013     1     1         2       11    1400    227     -9
## 2  2013     1     1         4       20    1416    227    -16
## 3  2013     1     1         2       33    1089    160    -31
## 4  2013     1     1        -1      -18    1576    183     17
## 5  2013     1     1        -6      -25     762    116     19
## 6  2013     1     1        -4       12     719    150    -16
## 7  2013     1     1        -5       19    1065    158    -24
## 8  2013     1     1        -3      -14     229     53     11
## 9  2013     1     1        -3       -8     944    140      5
## 10 2013     1     1        -2        8     733    138    -10
## # ... with 336,766 more rows, and 1 more variable: speed_miles <dbl>
```

No one knows what speed in miles is, let's fix that

Magic numbers. Great, every one loves them. They are evil.

```
KM_PER_MILE <- 1.61
```

```
mutate(flights_small,
  speed_km = (distance * KM_PER_MILE/air_time) * 60)
```

```
## # A tibble: 336,776 x 8
##   year month   day dep_delay arr_delay distance air_time speed_km
##   <int> <int> <int>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1  2013     1     1         2       11    1400    227    596.
## 2  2013     1     1         4       20    1416    227    603.
## 3  2013     1     1         2       33    1089    160    657.
## 4  2013     1     1        -1      -18    1576    183    832.
## 5  2013     1     1        -6      -25     762    116    635.
## 6  2013     1     1        -4       12     719    150    463.
## 7  2013     1     1        -5       19    1065    158    651.
## 8  2013     1     1        -3      -14     229     53    417.
## 9  2013     1     1        -3       -8     944    140    651.
## 10 2013     1     1        -2        8     733    138    513.
## # ... with 336,766 more rows
```

Even nicer is to create intermediate results for clarity

```
mutate(flights_small,
  distance_km = distance * KM_PER_MILE,
  air_time_hours = air_time / 60,
  speed_km = distance_km / air_time_hours
)
```

```
## # A tibble: 336,776 x 10
##   year month   day dep_delay arr_delay distance air_time distance_km
##   <int> <int> <int>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
```

```
## 1 2013 1 1 2 11 1400 227 2254
## 2 2013 1 1 4 20 1416 227 2280.
## 3 2013 1 1 2 33 1089 160 1753.
## 4 2013 1 1 -1 -18 1576 183 2537.
## 5 2013 1 1 -6 -25 762 116 1227.
## 6 2013 1 1 -4 12 719 150 1158.
## 7 2013 1 1 -5 19 1065 158 1715.
## 8 2013 1 1 -3 -14 229 53 369.
## 9 2013 1 1 -3 -8 944 140 1520.
## 10 2013 1 1 -2 8 733 138 1180.
## # ... with 336,766 more rows, and 2 more variables: air_time_hours <dbl>,
## # speed_km <dbl>
```

You cannot use all transformations inside mutate. It has to be vectorized: it takes a vector and returns a vector of the same length. The reason (I believe) is that the operation is done on the column as a whole. For this the operation needs to make sense for a whole column, not just for one number.

SOME VECTORIZED OPERATIONS

Standard arithmetic functions will work: $+$, $*$, etc

The time in dep_time is given by HHMM (How do I know this?)

```
transmute(flights,
  dep_time,
  dep_hour = dep_time %/% 100,
  dep_minutes = dep_time %% 100
)
```

```
## # A tibble: 336,776 x 3
##   dep_time dep_hour dep_minutes
##   <int>    <dbl>    <dbl>
## 1     517         5         17
## 2     533         5         33
## 3     542         5         42
## 4     544         5         44
## 5     554         5         54
## 6     554         5         54
## 7     555         5         55
## 8     557         5         57
## 9     557         5         57
## 10    558         5         58
## # ... with 336,766 more rows
```

log(), log2(), log10() work

How can you test whether something is vectorized?

```
(x <- c(0,1,2,3,4,5,6,7,8,9))

## [1] 0 1 2 3 4 5 6 7 8 9
```

```
(y <- 0:9)

## [1] 0 1 2 3 4 5 6 7 8 9
(z <- seq(0,9))

## [1] 0 1 2 3 4 5 6 7 8 9
(lag(y))

## [1] NA 0 1 2 3 4 5 6 7 8
(lag(lag(y)))

## [1] NA NA 0 1 2 3 4 5 6 7
(lead(y))

## [1] 1 2 3 4 5 6 7 8 9 NA
```

What do lag and lead do?

Some cumulative and aggregate functions

```
cumsum(x)

## [1] 0 1 3 6 10 15 21 28 36 45
cumprod(x)

## [1] 0 0 0 0 0 0 0 0 0 0
cumprod(lead(x))

## [1] 1 2 6 24 120 720 5040 40320 362880 NA
#?cummin
# ?cummax
cummean(x)

## [1] 0.0 0.5 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5
```

Logical operators work

```
x > 3

## [1] FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE
x > y

## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
x == y

## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
# What does the answer to this even mean?
x == c(2,4)

## [1] FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
x > c(2,4,6)
```

```
## Warning in x > c(2, 4, 6): longer object length is not a multiple of
## shorter object length
```

```
## [1] FALSE FALSE FALSE TRUE FALSE FALSE TRUE TRUE TRUE TRUE
```

Ranking functions

```
y <- c(10, 5, 6, 3, 7)
min_rank(y)
```

```
## [1] 5 2 3 1 4
```

Can you figure out from playing around with `min_rank()` how it works exactly?

So, what is not a vectorized operation?

```
c(2,4)^2 # This is vectorized
```

```
## [1] 4 16
```

```
kk <- function(x) { x[3]}
```

```
kk(1:5) # not vectorized
```

```
## [1] 3
```

```
mean(x)
```

```
## [1] 4.5
```

What happens when we try this on a dataframe

```
transmute(flights, delay = mean(arr_delay, na.rm = TRUE))
```

```
## # A tibble: 336,776 x 1
```

```
##   delay
```

```
##   <dbl>
```

```
## 1  6.90
```

```
## 2  6.90
```

```
## 3  6.90
```

```
## 4  6.90
```

```
## 5  6.90
```

```
## 6  6.90
```

```
## 7  6.90
```

```
## 8  6.90
```

```
## 9  6.90
```

```
## 10 6.90
```

```
## # ... with 336,766 more rows
```

```
transmute(flights, delay = kk(arr_delay))
```

```
## # A tibble: 336,776 x 1
##   delay
##   <dbl>
## 1    33
## 2    33
## 3    33
## 4    33
## 5    33
## 6    33
## 7    33
## 8    33
## 9    33
## 10   33
## # ... with 336,766 more rows
```

Exercise: Try out a few of the other commands in the chapter.

```
vars <- c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
```

```
row_number(vars)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11
```

```
dense_rank(vars)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11
```

```
percent_rank(vars)
```

```
## [1] 0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
```

```
cume_dist(vars)
```

```
## [1] 0.09090909 0.18181818 0.27272727 0.36363636 0.45454545 0.54545455
## [7] 0.63636364 0.72727273 0.81818182 0.90909091 1.00000000
```

Exercise: Create several ranges with the n:m notation, i.e. 2:4, 4:8, etc.

Try to find out whether you can also take negative ranges and descending.

Ans: Yes we can take negative numbers and descending as well based on the examples below.

```
vars_asc <- c(0:20)
```

```
vars_neg <- c(-5:-1)
```

```
vars_desc <- c(7:0)
```

Exercise: Read `?:` (the same as `help("?:")`)

Its the same as `help`.

Exercise: Use `slice()` to choose the first 10 rows of flights.

```
slice(flights, 1:10)
```

```
## # A tibble: 10 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int> <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
```

```
## 3 2013 1 1 542 540 2 923
## 4 2013 1 1 544 545 -1 1004
## 5 2013 1 1 554 600 -6 812
## 6 2013 1 1 554 558 -4 740
## 7 2013 1 1 555 600 -5 913
## 8 2013 1 1 557 600 -3 709
## 9 2013 1 1 557 600 -3 838
## 10 2013 1 1 558 600 -2 753
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

Do the following exercises from 5.5.2:

Exercise 1 Currently `dep_time` and `sched_dep_time` are convenient to look at, but hard to compute with because they're not really continuous numbers. Convert them to a more convenient representation of number of minutes since midnight.

```
#Convert into minutes from midnight
min_h <- 60
flights_updated <- flights %>%
  mutate(dep_time = (dep_time %/% 100)*min_h + dep_time %% 100,
         sched_dep_time = (sched_dep_time %/% 100)*min_h + sched_dep_time %% 100)
```

Exercise 2 Compare `air_time` with `arr_time - dep_time`. What do you expect to see? What do you see? What do you need to do to fix it?

Ans: Since the values of `arr_time` and `dep_time` are not continuous the subtraction leads to the wrong result and a value different from the pre-computed `air_time` value (which is in minutes). Both `arr_time` and `dep_time` need to be converted to minutes from midnight continuous values and then the `arr_time - dep_time` (journey_time in below solution) will be computed correctly.

```
flights %>% transmute(air_time, journey_time = arr_time - dep_time)

## # A tibble: 336,776 x 2
##   air_time journey_time
##   <dbl>      <int>
## 1     227         313
## 2     227         317
## 3     160         381
## 4     183         460
## 5     116         258
## 6     150         186
## 7     158         358
## 8      53         152
## 9     140         281
## 10    138         195
## # ... with 336,766 more rows
#journey_time is arr_time - dep_time
flights_updated <- flights %>%
  transmute(dep_time, arr_time, dep_time = (dep_time %/% 100)*min_h + dep_time %% 100,
         arr_time = (arr_time %/% 100)*min_h + arr_time %% 100, journey_time = arr_time - dep_time, air_time)

flights_updated %>% select(dep_time, arr_time, journey_time, air_time)

## # A tibble: 336,776 x 4
```

```
##   dep_time arr_time journey_time air_time
##   <dbl>   <dbl>       <dbl>   <dbl>
## 1     317     510         193     227
## 2     333     530         197     227
## 3     342     563         221     160
## 4     344     604         260     183
## 5     354     492         138     116
## 6     354     460         106     150
## 7     355     553         198     158
## 8     357     429          72      53
## 9     357     518         161     140
## 10    358     473         115     138
## # ... with 336,766 more rows
```

Exercise 4

Find the 10 most delayed flights using a ranking function. How do you want to handle ties? Carefully read the documentation for `min_rank()`.

Ans: The 10 most delayed flights can be found using `min_rank`. Ties are handled by giving the same rank to tied values and `min_rank` does it the same way.

```
del_flights <- flights %>% filter(min_rank(desc(dep_delay)) <= 10) %>% arrange(desc(dep_delay))
del_flights
```

```
## # A tibble: 10 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     9     641             900       1301    1242
## 2  2013     6    15    1432            1935       1137    1607
## 3  2013     1    10    1121            1635       1126    1239
## 4  2013     9    20    1139            1845       1014    1457
## 5  2013     7    22     845            1600       1005    1044
## 6  2013     4    10    1100            1900        960    1342
## 7  2013     3    17    2321             810        911     135
## 8  2013     6    27     959            1900        899    1236
## 9  2013     7    22    2257             759        898     121
## 10 2013    12     5     756            1700        896    1058
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

Hint: When you get stuck, try the following two strategies: 1. Take a single row, and work it out by hand 2. Create a variable `my_flights` which contains only a few rows (4 to 10). Work out a solution for `my_flights`, where you can check every step.

summarise()

```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 x 1
##   delay
##   <dbl>
## 1  12.6
```

How... useful. Might as well do


```
mean(flights$dep_delay, na.rm = TRUE)
```

```
## [1] 12.63907
```

\$ will give you that column. Quick way to choose columns.

```
mean(select(flights, dep_delay), na.rm = TRUE)
```

ERROR: argument is not numeric or logical: returning NA[1] NA

An error I made: I tried this: Huh? What's going on here?

```
flights$dep_delay  
select(flights, dep_delay)
```

I thought `select(flights, dep_delay)` was the same as `flights$dep_delay` Aha, we should have guessed, since `select` returns a *data frame*, but we want a column. A data frame of 1 column is not the same as a single column.

Still, summarise is way more interesting with its friend, `group_by`

```
by_day <- group_by(flights, year, month, day)  
by_day
```

```
## # A tibble: 336,776 x 19  
## # Groups:   year, month, day [365]  
##   year month   day dep_time sched_dep_time dep_delay arr_time  
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>  
## 1  2013     1     1     517             515         2     830  
## 2  2013     1     1     533             529         4     850  
## 3  2013     1     1     542             540         2     923  
## 4  2013     1     1     544             545        -1    1004  
## 5  2013     1     1     554             600        -6     812  
## 6  2013     1     1     554             558        -4     740  
## 7  2013     1     1     555             600        -5     913  
## 8  2013     1     1     557             600        -3     709  
## 9  2013     1     1     557             600        -3     838  
## 10 2013     1     1     558             600        -2     753  
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,  
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,  
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
## #   minute <dbl>, time_hour <dtm>
```

Looks distinctly the same

But it really isn't!

```
summarise(  
  group_by(flights, year, month, day),  
  delay = mean(dep_delay, na.rm = TRUE)  
)
```

```
## # A tibble: 365 x 4  
## # Groups:   year, month [12]  
##   year month   day delay  
##   <int> <int> <int> <dbl>  
## 1  2013     1     1 11.5  
## 2  2013     1     2 13.9  
## 3  2013     1     3 11.0
```

```
## 4 2013 1 4 8.95
## 5 2013 1 5 5.73
## 6 2013 1 6 7.15
## 7 2013 1 7 5.42
## 8 2013 1 8 2.55
## 9 2013 1 9 2.28
## 10 2013 1 10 2.84
## # ... with 355 more rows
```

5.6.1 Let's explore link between distance and average delay for every location. What that means is that we want to know the average delay for every destination. Then, once we have that, we want to see how the distance to this location is related to the delay to this location.

```
by_destination <- group_by(flights, dest)
delay <- summarise(by_destination,
                   delay = mean(arr_delay, na.rm = TRUE))
delay
```

```
## # A tibble: 105 x 2
##   dest   delay
##   <chr> <dbl>
## 1 ABQ    4.38
## 2 ACK    4.85
## 3 ALB   14.4
## 4 ANC   -2.5
## 5 ATL   11.3
## 6 AUS    6.02
## 7 AVL    8.00
## 8 BDL    7.05
## 9 BGR    8.03
## 10 BHM   16.9
## # ... with 95 more rows
```


OK, we need the distance too, or else there is not much to plot.

```
(delay <- summarise(by_destination,
                   delay = mean(arr_delay, na.rm = TRUE),
                   distance = mean(distance, na.rm = TRUE)))
```

```
## # A tibble: 105 x 3
##   dest   delay distance
##   <chr> <dbl>   <dbl>
## 1 ABQ    4.38   1826
## 2 ACK    4.85    199
## 3 ALB   14.4    143
## 4 ANC   -2.5   3370
## 5 ATL   11.3   757.
## 6 AUS    6.02  1514.
## 7 AVL    8.00   584.
## 8 BDL    7.05    116
## 9 BGR    8.03    378
## 10 BHM   16.9   866.
## # ... with 95 more rows
```

```
p <- ggplot(data = delay,
            mapping = aes(x = distance, y = delay))
p + geom_point() + geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
## Warning: Removed 1 rows containing non-finite values (stat_smooth).
## Warning: Removed 1 rows containing missing values (geom_point).
```



Improving the graph...

```
(delay <- summarise(by_destination,
  count = n(),
  delay = mean(arr_delay, na.rm = TRUE),
  distance = mean(distance, na.rm = TRUE)))
```


```
## # A tibble: 105 x 4
##   dest  count delay distance
##   <chr> <int> <dbl>    <dbl>
## 1 ABQ    254  4.38    1826
## 2 ACK    265  4.85     199
## 3 ALB   439 14.4     143
## 4 ANC      8 -2.5    3370
## 5 ATL  17215 11.3     757.
## 6 AUS  2439  6.02    1514.
## 7 AVL   275  8.00     584.
## 8 BDL   443  7.05     116
## 9 BGR   375  8.03     378
## 10 BHM   297 16.9     866.
## # ... with 95 more rows
```

```
p <- ggplot(data = delay,
            mapping = aes(x = distance, y = delay))
p + geom_point(mapping = aes(size = count), alpha = 0.2) +
  geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

```
## Warning: Removed 1 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```



$n()$ is a very special function $\#n()$

Finally...

Optional exercise as part of assignment 5 (somewhat harder): The above does not take into account

the number of flights per location. A location with 1 flight matters as much

for smoothing as a location with 300.

That is rarely what we want when smoothing globally. Read the following code,

to see if you understand how it works. Explain in your words in the .Rmd file.

Let's plot the original data, without first taking means by group

Woah, that looks different! (And ugly.)

So, not too misleading, but still...

END OF EXERCISE

Doing this with a pipe, and filtering out destinations with - less than 20 flights - to HNL (Honolulu), since it's by far the furthest Note: I am not a big fan of dropping things that 'look too different'. You should do such robustness checks, but you shouldn't start there.

```
delays <- flights %>%
  group_by(dest) %>%
  summarise(
    delay = mean(arr_delay, na.rm = TRUE),
    count = n(),
    distance = mean(distance, na.rm = TRUE)
  ) %>%
  filter( count > 20, dest != "HNL")
```

Exercise: Rewrite the above command without the pipe. Which one do you find easier to read?

Ans: Piping makes it much easier to write and much easier to read as well for me because I look at different steps without getting confused by extra information.

```
delays <- group_by(flights, dest)
delays <- summarise(delays, delay = mean(arr_delay, na.rm = TRUE),
  count = n(),
  distance = mean(distance, na.rm = TRUE))
filter(delays, count > 20, dest != 'HNL')
```

```
## # A tibble: 96 x 4
##   dest  delay count distance
##   <chr> <dbl> <int>    <dbl>
## 1 ABQ    4.38   254    1826
```

```
## 2 ACK      4.85    265    199
## 3 ALB     14.4    439    143
## 4 ATL     11.3  17215    757.
## 5 AUS      6.02   2439   1514.
## 6 AVL      8.00    275    584.
## 7 BDL      7.05    443    116
## 8 BGR      8.03    375    378
## 9 BHM     16.9     297    866.
## 10 BNA     11.8   6333    758.
## # ... with 86 more rows
```

5.6.2 Missing values

```
not_missing <- flights %>%
  filter(!is.na(dep_delay), !is.na(arr_delay))
```

Exercise: Does the above command also drop observations that miss only the arr_delay but have a dep_delay?

Ans: YES

Are there any observations in the dataset for which only dep_delay or arr_delay is missing, but not both?

Ans: No, because both of the commands below return 0 result set.

```
flights %>%
  filter(!is.na(dep_delay) & is.na(arr_delay)) %>% select(arr_delay, dep_delay)
```

```
## # A tibble: 0 x 2
## # ... with 2 variables: arr_delay <dbl>, dep_delay <dbl>
```


```
flights %>%
  filter(is.na(dep_delay) & !is.na(arr_delay)) %>% select(arr_delay, dep_delay)
```

```
## # A tibble: 0 x 2
## # ... with 2 variables: arr_delay <dbl>, dep_delay <dbl>
```

5.6.3 Counts


Average delay by airplane (identified by tailnum), plot density Start with freqpoly, then zoom in on that part of the graph that we are interested in..

```
not_missing %>%
  group_by(tailnum) %>%
  summarise(delay = mean(dep_delay)) %>%
  ggplot(mapping = aes(x = delay)) +
  geom_histogram(binwidth = 10)
```



Plot number of flights per airplane against delay

```
not_missing %>%  
  group_by(tailnum) %>%  
  summarise(  
    count = n(),  
    delay = mean(arr_delay)  
  ) %>%  
  ggplot(mapping = aes(x = delay, y = count)) +  
  geom_point(alpha = 0.1)
```




Since I need to filter the same thing, all the time just store in a variable. Delete other stuff.

```
not_missing_planes <- not_missing %>%
  group_by(tailnum) %>%
  summarise(
    count = n(),
    delay = mean(arr_delay),
    delay_median = median(arr_delay)
  )
```

Get the median delay for each ariplane


```
ggplot(data = not_missing_planes) +
  geom_histogram(mapping = aes(x = delay_median)) +
  geom_histogram(mapping = aes(x = delay), color = 'yellow', alpha = 0.3)
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
not_missing_planes %>%  
  filter(count > 5) %>%  
  ggplot(mapping = aes(x = delay)) +  
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



Filter the airplanes that fly rarely and pipe them into ggplot which gets plussed into geoms.

Try a few values for how many flights one should have done

Assignment 5:

1. Do the exercises in this script file and work through the examples we didn't cover in class. As usual, turn the script into an .Rmd file, knit it, upload the .html and .pdf.
2. Grade Assignment 4 of your peers.

4. Document at least 10 errors and warnings you actually hit during the week.

If you do *not* hit that many errors or receive such warnings, congratulations.

1. Row_number() ERROR: is.na() applied to non-(list or vector) of type 'closure' Error in x[!nas] : object of type 'closure' is not subsettable

Was passing a dataframe to is.na() instead of vector

2.dense_rank() ERROR: Error in unique.default(x) : unique() applies only to vectors

I was using an array of strings instead of integers

3. del_flights <- filter(flights, min_rank(dep_delay)) Error: Argument 2 filter condition does not evaluate to a logical vector - learned how to use min_rank and filter together

Was using a dataframe with one column instead of a vector

4. min and filter Error in min_rank(., desc(dep_delay)) : unused argument (desc(dep_delay)) -

Using `min_rank` inside filter - `min_rank` was not evaluating to a boolean/logical vector because I was missing out a comparison statement

5. Error in plotting Vienna Data against Hotels Data `ggplot(mapping = aes(x = viennaprice, y = hotels_dataprice))`

Both x and y need to be equal in length to make a scatter plot

6. Error in plotting side by side bar Could not figure out how to plot a side by side bar in ggplot - hit a number of errors and then gave up.

7. Had problems in leading ggarrange to combine multiple plots in one figure Error in `ggarrange()` : could not find function “ggarrange” Error in `library(ggpubr)` : there is no package called ‘ggpubr’

Fixed by installing the ggpubr package

8. Error in using hjust parameter to adjust the horizontal alignment of the labels of the combined plots Error: unexpected symbol in: “`ggarrange(vienna_plt, ams_plt, labels = c(“Vienna“, “Amsterdam“), hjust = c(-1, -1) ggarrange`”

Turns out we cannot adjust the horizontal alignment of both labels separately but it has to be a common value for both - need to pass an int instead of a list of ints

Pick one of the hotels graphs in Chapter 3, section 6, A1. Case study, finding a good deal among hotels. Replicate it – try it yourself for 10 minutes before you go looking at the code – and then make a variation of it.


```
vienna <- read_csv('hotels-vienna.csv')
```

```
## Parsed with column specification:
## cols(
##   .default = col_double(),
##   country = col_character(),
##   city_actual = col_character(),
##   center1label = col_character(),
##   center2label = col_character(),
##   neighbourhood = col_character(),
##   city = col_character(),
##   offer_cat = col_character(),
##   accommodation_type = col_character()
## )
```

```
## See spec(...) for full column specifications.
```

```
vienna %>% group_by(price) %>%
  summarise(frequency = n()) %>%
  ggplot(mapping = aes(x = price, y = frequency)) + geom_col() + xlim(0, 500)
```


```
## Warning: Removed 7 rows containing missing values (position_stack).
```



```
vienna %>% group_by(price) %>%  
  summarise(frequency = n()) %>%  
  ggplot(mapping = aes(x = price)) + geom_histogram(binwidth = 20, color = 'Grey') + xlim(0, 500)
```

```
## Warning: Removed 7 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 2 rows containing missing values (geom_bar).
```



6. Instead of using the Vienna data, use the data for another city (pick London if you don't want to choose). Do a basic data exploration, comparing the city to Vienna in terms of any variables you find interesting. Three plots maximum, don't spend more than 30 minutes on the analysis, before writing it down (if you are not doing this in parallel).

```
#two different sheets with price and distance (in _features sheet)
features <- read_csv('hotels-europe_features.csv')
```

```
## Parsed with column specification:
## cols(
##   hotel_id = col_double(),
##   city = col_character(),
##   distance = col_double(),
##   stars = col_double(),
##   rating = col_double(),
##   country = col_character(),
##   city_actual = col_character(),
##   rating_reviewcount = col_double(),
##   center1label = col_character(),
##   center2label = col_character(),
##   neighbourhood = col_character(),
##   ratingta = col_double(),
##   ratingta_count = col_double(),
##   distance_alter = col_double(),
##   accommodation_type = col_character()
## )
```

```

price <- read_csv('hotels-europe_price.csv')

## Parsed with column specification:
## cols(
##   hotel_id = col_double(),
##   price = col_double(),
##   offer = col_double(),
##   offer_cat = col_character(),
##   year = col_double(),
##   month = col_double(),
##   weekend = col_double(),
##   holiday = col_double(),
##   nnights = col_double(),
##   scarce_room = col_double()
## )

#merge the two tables
hotels_data <- merge(features, price, by = 'hotel_id')

ams_data <- filter(hotels_data, city == "Amsterdam" & stars == 3.0) #filtered by amsterdam

vienna_data <- vienna %>% filter(stars == 3.0)

v_stat <- vienna_data %>% group_by(city)

v_stat <- v_stat %>% summarise(mean = mean(price), median = median(price))
v_stat

## # A tibble: 1 x 3
##   city      mean median
##   <chr>    <dbl>  <dbl>
## 1 Vienna  106.    89

ams_stat <- ams_data %>% group_by(city)

ams_stat <- ams_stat %>% summarise(mean = mean(price), median = median(price))
ams_stat

## # A tibble: 1 x 3
##   city      mean median
##   <chr>    <dbl>  <dbl>
## 1 Amsterdam 275.   171

Compare the prices of 3 star hotels in Vienna and Amsterdam

vienna_plt <- vienna_data %>% ggplot(mapping = aes(x = price)) + geom_histogram(binwidth = 20, color = "red", fill = "white")

ams_plt <- ams_data %>% filter(stars == 3.0) %>% ggplot(mapping = aes(x = price)) + geom_histogram(binwidth = 20, color = "blue", fill = "white")

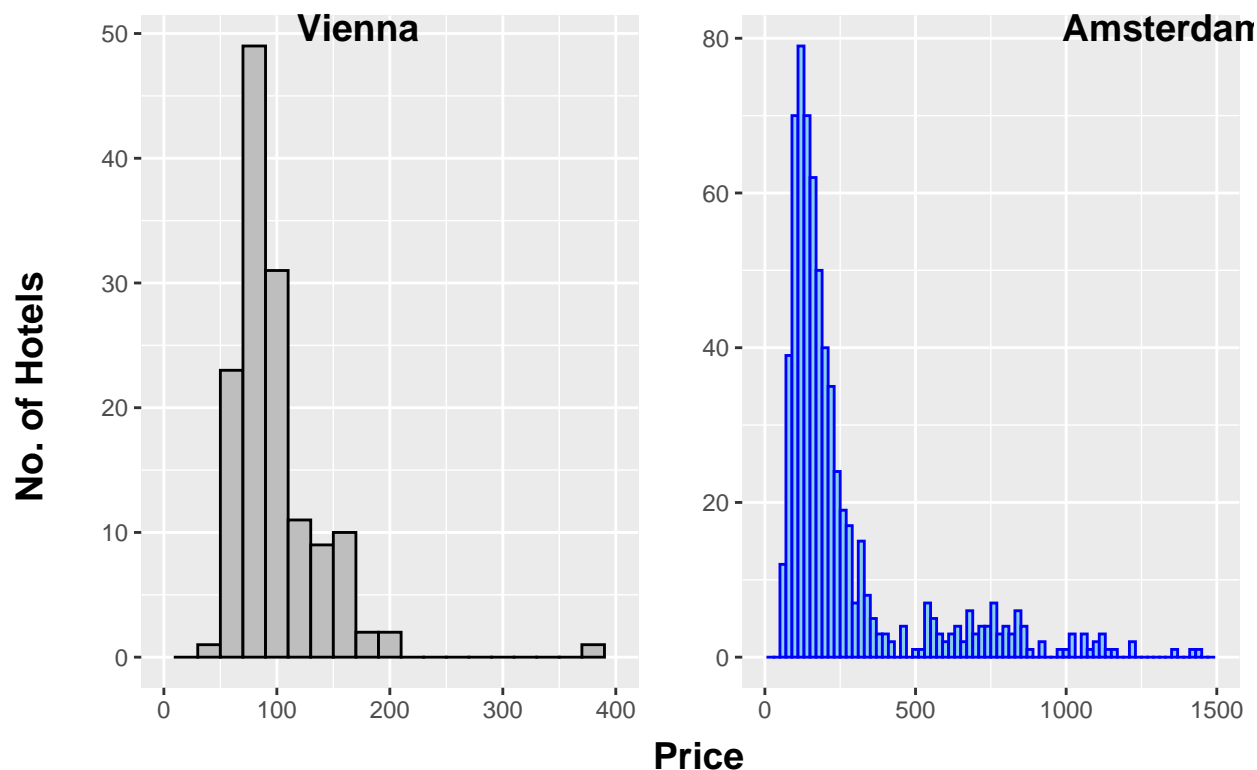
combined_figure <- ggarrange(vienna_plt, ams_plt, labels = c("Vienna", "Amsterdam"), hjust = -2)

annotate_figure(combined_figure,
  top = text_grob("Price Comparison for Three Star Hotels", color = "Grey", face = "bold", size = 14, hjust = 0, x = 0.5),
  bottom = text_grob("Price", vjust = -1, hjust = 0, x = 0.5, face = "bold", size = 14),
  left = text_grob("No. of Hotels", rot = 90, size = 14, face = "bold"),

```

)

Price Comparison for Three Star Hotels



```
ams_data <- filter(hotels_data, city == "Amsterdam") #filtered by amsterdam
```

```
vienna_data <- vienna
```

```
v_stat <- vienna_data %>% group_by(city)
```

```
v_stat <- v_stat %>% summarise(mean = mean(price), median = median(price))
```

```
v_stat
```

```
## # A tibble: 1 x 3
##   city    mean median
##   <chr>  <dbl>  <dbl>
## 1 Vienna  131.   110.
```

```
ams_stat <- ams_data %>% group_by(city)
```

```
ams_stat <- ams_stat %>% summarise(mean = mean(price), median = median(price))
```

```
ams_stat
```

```
## # A tibble: 1 x 3
##   city    mean median
##   <chr>  <dbl>  <dbl>
## 1 Amsterdam 318.   196
```

Compare the overall spread of prices of hotels in Vienna and Amsterdam

```

vienna_plt <- vienna_data %>% ggplot(mapping = aes(x = price)) + geom_histogram(binwidth = 20, color = "grey", fill = "white")

ams_plt <- ams_data %>% filter(stars == 3.0) %>% ggplot(mapping = aes(x = price)) + geom_histogram(binwidth = 20, color = "blue", fill = "white",
  labs(x = "", y = ""))

combined_figure <- ggarrange(vienna_plt, ams_plt, labels = c("Vienna", "Amsterdam"), hjust = -2)

annotate_figure(combined_figure,
  top = text_grob("Price Comparison for all hotels in Amsterdam and Vienna",
    color = "grey", face = "bold", size = 14),
  bottom = text_grob("Price", vjust = -1,
    hjust = 0, x = 0.5, face = "bold", size = 14),
  left = text_grob("No. of Hotels", rot = 90, size = 14, face = "bold"),
)

```

